Interpreting Non-Coding Genetic Variation with Machine Learning: Review

1 Background

Most genetic variants linked to disease lie outside protein-coding regions. These non-coding variants can change how genes are regulated by altering transcription factor binding, promoter or enhancer activity, or splicing signals. Experimentally testing every variant is impractical because there are simply too many possible variiants, so computational models that predict which variants are likely functional and in which cell types they matter are essential for research and precision medicine.

2 Papers

The two models I will be reviewing in this paper are Bio-DTA and DeepVRegulome. Both models aim to explore machine learning approaches to exploring the space of non-coding variants.

DeepVRegulome aims to find and prioritize short variants in regulatory DNA that disrupt transcription factor binding sites or splice motifs. The authors fine-tune many DNABERT transformer classifiers on ENCODE and GENCODE regulatory annotations, score variants from whole genome sequences of glioblastoma patients and test whether prioritized variants are associated with patient survival.

Bio-DTA aims to connect DNA sequence changes to predict gene expression effects at single-cell resolution. The core idea is to map DNA-derived embeddings of genes onto the token embedding space of a pre-trained single-cell expression model allowing for in-silico sequence perturbations can be assessed in a cell-type to cell-type context.

These two studies are complementary in the sense that DeepVRegulome answers which variants are more likely to disrupt regulatory sequence features, while Bio-DTA explores how those disruptions might alter gene expression in a specific cell-type.

2.1 Machine Learning Methods

DeepVRegulome uses DNABERT-style transformers. A transformer is a neural network that models sequences by learning relationships among positions through a self-attention mechanism which allows it to capture both short motifs and longer context. The authors fine-tune many transformer models on regulatory windows from ENCODE and GENCODE so each model specializes on a transcription factor or splice motif. To score a variant, the model compares its output on the reference sequence to the alternate sequence, with large changes indicating likely disruption.

Bio-DTA combines DNA sequence embeddings with a pre-trained single-cell foundation model. A single-cell foundation model is a large neural network trained on millions of single-cell RNA profiles to learn typical gene expression programs across cell types. Bio-DTA uses enformer-like sequence embeddings for each gene and a learned adapter that maps those embeddings into the token space of the single-cell model. Changing the input embedding for a gene simulates a DNA perturbation, and the downstream changes in the single-cell model indicate predicted expression effects in a chosen cell type. This maps genotype to cell-type-aware expression predictions.

2.2 Outcomes

The fine-tuned models for DeepVRegulome show strong performance on known regulatory sites. The splice models report high accuracy and F1 scores near the mid 90s, and top TF models have AUCs above 0.95 in validation. Applied to glioblastoma whole-genome data, the pipeline produced a catalog of thousands of candidate regulatory variants and identified sets of mutations and regions associated with patient survival, suggesting clinical relevance.

Bio-DTA demonstrated in-silico perturbations, such as mutating the GATA4 promoter and observing predicted expression shifts of known targets in fetal cardiomyocytes but not in unrelated cell types. Quantitatively, the approach increased precision and recall for recovering expected targets compared to baselines in the presented case studies, and embedding shifts were statistically significant for direct targets. However, large-scale ground-truth benchmarks linking precise sequence changes to cell-type expression remain limited, so their evaluations are currently case-study focused.

2.3 Strengths and Limitations

DeepVRegulome is scalable and interpretable at the motif level and links predictions to clinical data, but it focuses on local sequence windows and lacks explicit cell-type context (the context that Bio-DTA's method may provide). Bio-DTA provides cell-type-specific functional hypotheses and a flexible adapter that bridges sequence and expression spaces, but it depends on the coverage and quality of single-cell training data and needs broader experimental benchmarking.

3 Conclusion

In conclusion, DeepVRegulome and Bio-DTA address complementary questions in non-coding variant interpretation. One prioritizes and interprets likely regulatory disruptions at scale and the other places those disruptions in a realistic cellular context by predicting expression consequences. A future direction in integrating the two methods, genome-scale variant scoring with cell-type-aware expression simulation, is a promising direction for producing testable, biologically relevant hypotheses, but currently both methods require more extensive validation and broader data coverage to reach routine use in clinical and experimental workflows. Overall, these papers advance the field by showing that deep learning can translate raw genomic sequence variation into meaningful biological predictions while circumventing extensive experimental work.

4 References

Sathian, R., Chao, M., Surana, P., Liu, H., Yurovsky, A., & Davuluri, R., V. (2025). DeepVRegulome: DNABERT-based deep-learning framework for predicting the functional impact of short genomic variants on the human regulome. arXiv (Cornell University). https://doi.org/10.48550/arxiv.2511.09026

Solaguren-Beascoa, A., Neilson, G., Muhammed, L., & Cakiroglu, S. A. (2025). Multi-modal single-cell foundation models via dynamic token adaptation. arXiv (Cornell University). https://doi.org/10.48550/arxiv.2504.13049 Paper Links - DeepVRegulome - https://arxiv.org/pdf/2511.09026 Multi-Modal Single-Cell Foundation Models via Dynamic Token Adaptation - https://arxiv.org/pdf/2504.13049